

Dynamic Alert Configuration System

Feasibility & Approach Document

Version: 1.0

Date: January 2026

Status: Feasibility Analysis for POC

Executive Summary

This document outlines a feasible approach to enable dynamic, customer-specific alert configuration without manual YAML editing or system downtime. The solution allows platform administrators to configure, modify, and manage alerts for multiple customers through a REST API and web interface while maintaining system stability and isolation.

Key Outcome: Customers can have fully customized alerts, notification preferences, and routing rules that can be updated in real-time without affecting other customers or requiring system restarts.

Problem Statement

Current Limitations

1. Manual Configuration Required

- Alert rules exist in static YAML files
- Any change requires manual file editing
- Risk of syntax errors and misconfigurations

2. No Multi-Tenancy Support

- Cannot isolate alerts per customer
- Single configuration affects all customers
- No customer-specific customization

3. Operational Risk

- Changes require Prometheus restart (downtime)
- Editing one customer's alerts can break others
- No validation before applying changes

4. Scalability Concerns

- Manual process doesn't scale beyond 5-10 customers
- Configuration drift across environments
- No audit trail of changes

Business Requirements

From User Story 1, the system must support:

- Configurable alert definitions (application & infrastructure)
- Per-customer thresholds and conditions
- Dynamic enable/disable without downtime
- Custom notification templates per customer
- Isolation between customers
- Audit trail of all changes

Proposed Solution

High-Level Architecture





Core Components

1. PostgreSQL Database

Purpose: Central source of truth for all alert configurations

Stores:

- Alert definitions with full PromQL expressions
- Customer-specific thresholds and parameters
- Notification receiver configurations
- Email/Slack/PagerDuty templates per customer
- Routing rules per customer
- Complete audit log of all changes

Benefits:

- Transactional consistency
- ACID guarantees
- Queryable audit trail
- Backup/restore capabilities

2. Alert Management API

Purpose: Provides programmatic interface for alert configuration

Capabilities:

- Create/Read/Update/Delete alert definitions
- Enable/disable alerts per customer
- Manage notification receivers

- Configure routing rules
- Validate PromQL expressions before saving
- Bulk import/export for migrations

Benefits:

- Version-controlled configuration
- API-first design for automation
- Input validation before persistence
- Role-based access control ready

3. Configuration Sync Engine

Purpose: Bridges database and Prometheus/Alertmanager

Workflow:

1. Triggered when alert configuration changes
2. Fetches all enabled alerts from database
3. Groups alerts by category (application/infrastructure)
4. Generates complete alerts.yml file in memory
5. Validates using Prometheus promtool
6. Writes atomically to filesystem
7. Triggers Prometheus hot reload via HTTP API
8. Repeats for Alertmanager configuration

Benefits:

- Automatic synchronization
- Validation before apply
- Atomic updates (all-or-nothing)
- Non-blocking operation

4. Hot Reload Mechanism

Purpose: Apply configuration changes without downtime

How it works:

- Prometheus provides `(/-reload)` HTTP endpoint

- Accepts POST request to reload configuration
- Re-reads YAML files without restart
- Typically completes in 1-5 seconds
- Existing queries continue during reload

Benefits:

- Zero downtime
 - No service interruption
 - Fast propagation (seconds, not minutes)
 - Built-in Prometheus capability
-

Technical Approach

Step 1: Database Schema Design

Alert Definitions Table

- Stores complete PromQL expressions (not decomposed)
- Includes customer_id for isolation
- Contains all metadata (severity, urgency, category)
- Tracks enabled/disabled status
- Records creation/modification timestamps

Notification Receivers Table

- Stores email addresses per customer
- Contains customizable email subject templates
- Contains customizable email body templates (HTML)
- Supports multiple channel types (email, Slack, PagerDuty)
- Links to specific customers

Routing Rules Table

- Defines which alerts go to which receivers
- Supports severity-based routing (critical vs warning)
- Supports category-based routing (application vs infrastructure)

- Priority-ordered for complex routing logic
- Customer-specific configuration

Audit Log Table

- Records every create/update/delete operation
- Captures before/after values
- Tracks who made the change and when
- Enables compliance and troubleshooting

Step 2: API Implementation

REST Endpoints:

- `POST /customers/{id}/alerts` - Create new alert
- `PUT /customers/{id}/alerts/{alertId}` - Update existing alert
- `PATCH /customers/{id}/alerts/{alertId}/enabled` - Toggle on/off
- `DELETE /customers/{id}/alerts/{alertId}` - Remove alert
- `GET /customers/{id}/alerts` - List all alerts
- `POST /customers/{id}/receivers` - Configure notifications
- `POST /customers/{id}/routing` - Configure routing

Request Flow:

1. API receives request with alert configuration
2. Validates PromQL syntax using Prometheus promtool
3. Validates customer_id is present in expression
4. Checks for duplicate alert names
5. Persists to database
6. Triggers async sync job
7. Returns immediately (non-blocking)
8. Sync completes in background (1-5 seconds)

Validation:

- PromQL expression syntax validation
- Required field validation
- Duplicate name detection

- Customer ID isolation enforcement
- Threshold value range checking

Step 3: YAML Generation Process

Single File Approach: All customers' alerts stored in one `alerts.yml` file, isolated by `customer_id` label.

Generation Logic:

1. Query database for all enabled alerts
2. Group by category (application_alerts, infrastructure_alerts)
3. For each alert:
 - Use PromQL expression exactly as stored
 - Add `customer_id` label
 - Add severity, urgency, category labels
 - Add annotations (summary, description, impact, action)
4. Build complete YAML structure in memory
5. Validate using promtool before writing
6. Write to temporary file
7. Atomically rename to `alerts.yml`

Alertmanager Configuration:

1. Query database for all customer receivers and routing rules
2. Build routing tree:
 - Root route groups by `customer_id`
 - Each customer gets sub-routes for severity/category
 - Each route points to customer-specific receiver
3. Generate receiver configurations:
 - Email receivers with custom templates
 - Slack receivers with webhooks
 - PagerDuty receivers with integration keys
4. Write to temporary file
5. Atomically rename to `alertmanager.yml`

Atomic Updates:

- Always write to temporary file first
- Validate before committing
- Use atomic rename operation
- Prevents partial/corrupted configurations

Step 4: Hot Reload Execution

Prometheus Reload:

1. Send HTTP POST to `(http://prometheus:9090/-/reload)`
2. Prometheus re-reads configuration files
3. Compiles new alert rules
4. Activates new rules within 1-5 seconds
5. Existing metrics collection continues uninterrupted
6. Query API remains available during reload

Alertmanager Reload:

1. Alertmanager watches config file for changes
2. Automatically reloads on file modification
3. Or send SIGHUP signal for immediate reload
4. Routing updates take effect within seconds
5. In-flight alerts continue processing

Error Handling:

- If validation fails, keep old configuration
- Log errors for troubleshooting
- Notify administrators of sync failures
- Implement retry logic with exponential backoff

Step 5: Customer Isolation

Label-Based Isolation: Every metric must include `(customer_id)` label:

```
http_request_duration_seconds{customer_id="customer-abc", service="nmt"}
```

PromQL Expression Requirements:

- All expressions must filter by customer_id
- API validates customer_id is present
- Prevents cross-customer data leakage
- Enables per-customer alerting

Routing Isolation:

- Alertmanager routes first by customer_id
 - Each customer's alerts go to their receivers only
 - Email templates are customer-specific
 - Complete notification isolation
-

Feasibility Analysis

Technical Feasibility:  HIGH

Evidence:

1. Prometheus Hot Reload is Standard

- Built-in capability since Prometheus 2.0
- Used in production by thousands of companies
- Well-documented and stable

2. YAML File Size is Manageable

- Average alert rule: ~500 bytes
- 100 customers × 20 alerts = 1 MB file
- Prometheus handles 10,000+ rules easily
- Reload time: 2-5 seconds for 10,000 rules

3. Database-Driven Configuration is Proven

- Standard pattern in SaaS platforms
- PostgreSQL handles this workload easily
- Transactional guarantees ensure consistency

4. API-First Approach is Industry Standard

- Similar to how Datadog, New Relic work
- RESTful design is well-understood

- Easy to integrate with UI or automation

Potential Challenges:

- Learning curve for PromQL expressions
- Need to ensure customer_id in all metrics
- Monitoring the sync process itself

Mitigations:

- Provide PromQL templates for common patterns
- API validation enforces customer_id requirement
- Health checks and logging for sync process

Operational Feasibility: HIGH

Benefits:

1. Zero Downtime Updates

- Hot reload means no service interruption
- Changes apply in seconds
- No maintenance windows required

2. Customer Isolation

- One customer's changes don't affect others
- Database transactions ensure consistency
- Label-based filtering prevents data leakage

3. Audit Trail

- Every change is logged with timestamp and user
- Can roll back to previous configurations
- Meets compliance requirements

4. Validation Before Apply

- Invalid configurations are rejected
- Production system protected from bad configs
- Clear error messages for troubleshooting

Operational Requirements:

- Monitoring of sync process (alerts on sync failures)
- Backup strategy for database
- Documentation for support team

Scalability:  **HIGH**

Performance Characteristics:

Customers	Alerts/Customer	Total Alerts	File Size	Reload Time	Sync Time
10	20	200	~100 KB	<1 sec	<1 sec
50	20	1,000	~500 KB	1-2 sec	1-2 sec
100	20	2,000	~1 MB	2-3 sec	2-3 sec
500	20	10,000	~5 MB	3-5 sec	3-5 sec

When to Consider Alternatives:

- 1,000+ customers: Evaluate Grafana Mimir (multi-tenant Prometheus with Rules API)
- 50,000+ rules: Consider sharding or recording rules
- 100+ changes per minute: Implement batching

For POC Scale (10-100 customers):

- Single file approach is optimal
- Hot reload is fast enough
- Database is not bottleneck
- No architectural changes needed

Implementation Complexity:  **MEDIUM**

Components to Build:

1. Database schema (1-2 days)
2. REST API (3-5 days)
3. Sync engine (3-5 days)
4. Validation logic (2-3 days)
5. Web UI (5-7 days, optional for POC)